



Université de Thiès  
UFR SES  
Master en Science des Données et Applications / Option IA  
Deep Learning (CM : Pr M.BOUSSO – M.LANDU)

## **Rapport Projet Régression Linéaire Multiple**

### **# Exercice 1: Régression Linéaire**

#### **Importation des librairies a utiliser**

```
import import_ipynb

# used for manipulating directory paths

import os


# Scientific and vector computation for python

import numpy as np


# Plotting library

from matplotlib import pyplot as plt

from mpl_toolkits.mplot3d import Axes3D # needed to plot 3-D surfaces


# library written for this exercise providing additional functions for assignment
submission, and others


import utils


# tells matplotlib to embed plots within the notebook

%matplotlib inline
```

## Création d'une Simple python and numpy function

```
def warmUpExercise(n):
```

```
    A = np.eye(n)
```

```
    return A
```

### Appel de la fonction

```
A=warmUpExercise(5)
```

```
print("Matrice d'identité d'ordre 5 ")
```

```
print(A)
```

### Résultat :

```
Matrice d'identité d'ordre 5
[[1.  0.  0.  0.  0.]
 [0.  1.  0.  0.  0.]
 [0.  0.  1.  0.  0.]
 [0.  0.  0.  1.  0.]
 [0.  0.  0.  0.  1.]]
```

---

## 2 Régression linéaire à une seule variable

### Code

```
# Read comma separated data
```

```
data = np.loadtxt(os.path.join('Data', 'ex1data1.txt'),delimiter=',')
```

```
X = data[:,0]
```

```
y = data[:,1]
```

```
m = X.size # number of training examples
```

```
# print out some data points
```

```
print('{:>8s}{:>10s}'.format('Pop(X)', 'profit(y)'))
```

```
print('-'*26)
```

```
for i in range(10):
```

```
print('{:8.0f}{:10.0f}'.format(X[i], y[i]))
```

## **Résultat**

Pop(X)	profit(y)
6	18
6	9
9	14
7	12
6	7
8	12
7	4
9	12
6	7
5	4

## **2.1 Creation fonction Plotting the Data**

```
def plotData(x, y,xlabel,ylabel)

    from matplotlib import pyplot as plt

    fig = plt.figure() # open a new figure

    plt.plot(x, y, 'ro', ms=10, mec='k')


    plt.ylabel('Profit in $10,000')

    plt.xlabel('Population of City in 10,000s')
```

### **Appel de la fonction**

```
plotData(X, y,'Population of City in 10,000s','Profit in $10,000')
```

### **Resulat :**



### **2.2.3 Calcul de la fonction coût ( $\theta$ )**

```
def computeCost(X, y, theta):  
    # initialize some useful values  
    m=len(y)  
    # You need to return the following variables correctly  
    J = 1/(2*m)*np.sum((model(X,theta)-y)**2)  
    return J
```

#### **Appel de fonction**

```
J = computeCost(X, y, theta=np.array([0.0, 0.0]))  
print('With theta = [0, 0] \nCost computed = %.2f' % J)  
print('Expected cost value (approximately) 32.07\n')  
# further testing of the cost function  
J = computeCost(X, y, theta=np.array([-1, 2]))  
print('With theta = [-1, 2]\nCost computed = %.2f' % J)  
print('Expected cost value (approximately) 54.24')
```

#### **Résultat**

```
With theta = [0, 0]  
Cost computed = 32.07  
Expected cost value (approximately) 32.07  
  
With theta = [-1, 2]  
Cost computed = 54.24  
Expected cost value (approximately) 54.24
```

### **2.2.4 Gradient descent**

```
def grad(X, y, theta):
```

```

m = len(y)

return 1/m * X.T.dot(model(X, theta) - y)

def gradientDescent(X, y, theta, alpha, num_iters):

    theta = theta.copy()

    J_history = [] # Use a python list to save cost in every iteration

    J_history=np.zeros(num_iters)

    for i in range (0,num_iters):

        theta=theta-alpha*grad(X,y,theta)

        J_history[i]=computeCost(X,y,theta)

    # save the cost J in every iteration

    return theta, J_history

```

### **Appel de fonction**

```

# initialize fitting parameters

theta = np.zeros(2)


# some gradient descent settings

iterations = 1500

alpha = 0.01


theta, J_history = gradientDescent(X ,y, theta, alpha, iterations)

print('Theta found by gradient descent: {:.4f}, {:.4f}'.format(*theta))

print('Expected theta values (approximately): [-3.6303, 1.1664]')

```

### **Résultat**

```

Theta found by gradient descent: -3.6303, 1.1664
Expected theta values (approximately): [-3.6303, 1.1664]

```

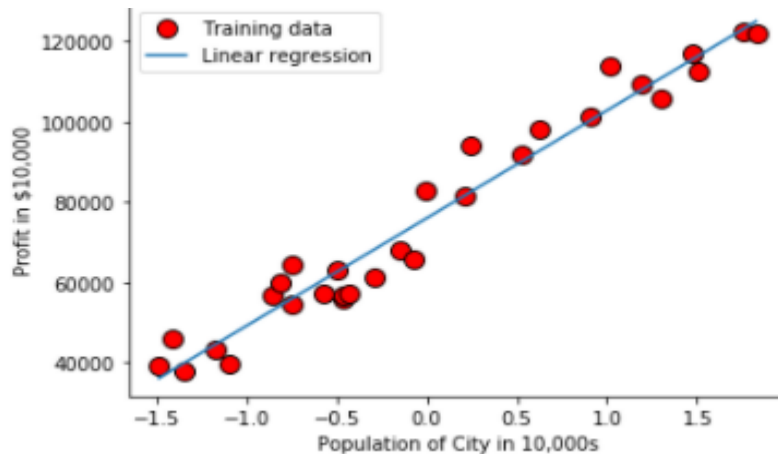
---

## # plot the linear fit

```
plotData(X[:, 1], y, 'Population of City in 10,000s', 'Profit in $10000')
```

```
plt.plot(X[:, 1], np.dot(X, theta), '-')
```

```
plt.legend(['Training data', 'Linear regression'])
```



---

## # Predict values for population sizes of 35,000 and 70,000

```
predict1 = np.dot([1, 3.5], theta)
```

```
print('For population = 35,000, we predict a profit of {:.2f}\n'.format(predict1*10000))
```

```
predict2 = np.dot([1, 7], theta)
```

```
print('For population = 70,000, we predict a profit of {:.2f}\n'.format(predict2*10000))
```

---

```
For population = 35,000, we predict a profit of 1698657108.94
```

```
For population = 70,000, we predict a profit of 2637284433.48
```

## Exercice: Choisir 3 points dans le contour et représenter les droites qui leur correspondent dans le nuage de points

### Code

```
from matplotlib import pyplot
```

```
# plot the linear fit


plotData(X[:, 1], y, 'Population of City in 10,000s', "Profit in $10000")

for i in range(3):

    theta=[np.random.choice(theta0_vals), np.random.choice(theta1_vals)]

    pyplot.plot(X[:, 1], np.dot(X, theta), '-')

plt.legend(['Training data', 'Linear regression']);
```



## **Exercice optionnel:**

On Reprend les mêmes cellules en changeant de cost function, prenez la moyenne des valeurs absolues des erreurs ( MAE).

### **Code**

```
def model(X, theta):

    return X.dot(theta)

def computeCostAbsolute(X, y, theta):

    J =(1/m)*np.sum(np.abs(model(X,theta)-y))

    return J
```

### **Appel Fonction**

```
J = computeCostAbsolute(X, y, theta=np.array([0.0, 0.0]))

print('With theta = [0, 0] \nCost computed = %.2f' % J)
```

```

print('Expected cost value (approximately) 76003.00\n')

# further testing of the cost function

J = computeCostAbsolute(X, y, theta=np.array([-1, 2]))

print('\n\nWith theta = [-1, 2]\nCost computed = %.2f' % J)

print('Expected cost value (approximately) 76004.00\n')

```

## **Résultat**

```

Entrée [145]: J = computeCostAbsolute(X, y, theta=np.array([0.0, 0.0]))
print('With theta = [0, 0] \nCost computed = %.2f' % J)
print('Expected cost value (approximately) 76003.00\n')
# further testing of the cost function
J = computeCostAbsolute(X, y, theta=np.array([-1, 2]))
print('\n\nWith theta = [-1, 2]\nCost computed = %.2f' % J)
print('Expected cost value (approximately) 76004.00\n')

With theta = [0, 0]
Cost computed = 76003.00
Expected cost value (approximately) 76003.00

With theta = [-1, 2]
Cost computed = 76004.00
Expected cost value (approximately) 76004.00

```

## **Régression linéaire avec plusieurs variables**

Dans cette partie, nous allons implémenter une régression linéaire avec plusieurs variables pour prédire les prix des maisons.

### **Feature Normalization**

```

# Load data

data = np.loadtxt(os.path.join('Data', 'ex1data2.txt'), delimiter=',')

X = data[:, :2]

y = data[:, 2]

m = y.size

```



```
# print out some data points

print('{:>8s}{:>8s}{:>10s}'.format('X[:,0]', 'X[:, 1]', 'y'))

print('-'*26)

for i in range(10):

    print('{:8.0f}{:8.0f}{:10.0f}'.format(X[i, 0], X[i, 1], y[i]))
```

```
# Load data
data = np.loadtxt(os.path.join('Data', 'ex1data2.txt'), delimiter=',')
X = data[:, :2]
y = data[:, 2]
m = y.size

# print out some data points
print('{:>8s}{:>8s}{:>10s}'.format('X[:,0]', 'X[:, 1]', 'y'))
print('-'*26)
for i in range(10):
    print('{:8.0f}{:8.0f}{:10.0f}'.format(X[i, 0], X[i, 1], y[i]))
```

X[:,0]	X[:, 1]	y
2104	3	399900
1600	3	329900
2400	3	369000
1416	2	232000
3000	4	539900
1985	4	299900
1534	3	314900
1427	3	198999
1380	3	212000
1494	3	242500

On complète le code dans la fonction `featureNormalize`:

```
def featureNormalize(X):
```

```
    # You need to set these values correctly
    X_norm = X.copy()
    mu = np.zeros(X.shape[1])
```

```

sigma = np.zeros(X.shape[1])
n = X_norm.shape[1]

for i in range(n):
    mu[i]=np.mean(X[:,i])
    sigma[i]=np.std(X[:,i])
    for j in range(len(X_norm)):
        X_norm[j,i]=(X_norm[j,i]-mu[i])/sigma[i]

return X_norm, mu, sigma

```

**# call featureNormalize on the loaded data**

```
X_norm, mu, sigma = featureNormalize(X)
```

```

print('Computed mean:', mu)
print('Computed standard deviation:', sigma)

```

**# Computed mean: [2000.68085106 3.17021277]**

**# Computed standard deviation: [7.86202619e+02 7.52842809e-01]**

**Resultat :**

```

Computed mean: [2000.68085106 3.17021277]
Computed standard deviation: [7.86202619e+02 7.52842809e-01]

```

## **Gradient Descent**

On complète le code des fonctions computeCostMulti et gradientDescentMulti pour implémenter la fonction de coût et la descente de gradient pour la régression linéaire avec plusieurs variables

```
def model(X, theta):
```

```
    return X.dot(theta)
```

```
def computeCostMulti(X, y, theta):
```

```
    """
```

```
    Compute cost for linear regression with multiple variables.
```

```
    Computes the cost of using theta as the parameter for linear regression to fit the data points in X
    and y.
```

## Parameters

-----

X : array\_like

The dataset of shape (m x n+1).

y : array\_like

A vector of shape (m, ) for the values at a given data point.

theta : array\_like

The linear regression parameters. A vector of shape (n+1, )

## Returns

-----

J : float

The value of the cost function.

## Instructions

-----

Compute the cost of a particular choice of theta. You should set J to the cost.

"""

# Initialize some useful values

m = y.shape[0] # number of training examples

# You need to return the following variable correctly

J = 0

J = (0.5/m) \* np.sum((model(X, theta) - y)\*\*2)

return J

def grad(X, y, theta):

```
m = len(y)
```

```
return 1/m * X.T.dot(model(X, theta) - y)
```

```
def gradientDescentMulti(X, y, theta, alpha, num_iters):
```

```
    """
```

Performs gradient descent to learn theta.

Updates theta by taking num\_iters gradient steps with learning rate alpha.

Parameters

-----

X : array\_like

The dataset of shape (m x n+1).

y : array\_like

A vector of shape (m, ) for the values at a given data point.

theta : array\_like

The linear regression parameters. A vector of shape (n+1, )

alpha : float

The learning rate for gradient descent.

num\_iters : int

The number of iterations to run gradient descent.

## Returns

-----

`theta : array_like`

The learned linear regression parameters. A vector of shape  $(n+1, )$ .

`J_history : list`

A python list for the values of the cost function after each iteration.

## Instructions

-----

Perform a single gradient step on the parameter vector `theta`.

While debugging, it can be useful to print out the values of the cost function (`computeCost`) and gradient here.

"""

# Initialize some useful values

`m = y.shape[0]` # number of training examples

# make a copy of theta, which will be updated by gradient descent

`theta = theta.copy()`

`J_history = []`

`J_history=np.zeros(num_iters)`

```

for i in range(0,num_iters):

    theta=theta-alpha*grad(X,y,theta)

    J_history[i]=computeCostMulti(X,y,theta)

return theta, J_history

```

## # initialize fitting parameters

```
theta = np.zeros(3)
```

## # some gradient descent settings

```
iterations = 1500
```

```
alpha = 0.01
```

```
theta, J_history = gradientDescentMulti(X ,y, theta, alpha, iterations)
```

```
print('Theta found by gradient descent: {:.4f}, {:.4f}, {:.4f}'.format(*theta))
```

## Résultat

```

]: # initialize fitting parameters
   theta = np.zeros(3)

   # some gradient descent settings
   iterations = 1500
   alpha = 0.01

   theta, J_history = gradientDescentMulti(X ,y, theta, alpha, iterations)
   print('Theta found by gradient descent: {:.4f}, {:.4f}, {:.4f}'.format(*theta
   #J_history

```

```
Theta found by gradient descent: 340412.5630, 109370.0567, -6500.6151
```

---

## **Exercise: Selection de taux d'apprentissage**

Dans cette partie de l'exercice, on va essayer différents taux d'apprentissage pour l'ensemble de données et trouver un taux d'apprentissage qui converge rapidement. On va aussi modifier le taux d'apprentissage en modifiant le code suivant et en modifiant la partie du code qui définit le taux d'apprentissage.

"""

Instructions

-----

We have provided you with the following starter code that runs gradient descent with a particular learning rate (alpha).

Your task is to first make sure that your functions - ``computeCost`` and ``gradientDescent`` already work with this starter code and support multiple variables.

After that, try running gradient descent with different values of alpha and see which one gives you the best result.

Finally, you should complete the code at the end to predict the price of a 1650 sq-ft, 3 br house.

Hint

----

At prediction, make sure you do the same feature normalization.

"""

**# Choose some alpha value - change this**

alpha = 0.1

num\_iters = 400

**# init theta and run gradient descent**

theta = np.zeros(3)

theta, J\_history = gradientDescentMulti(X, y, theta, alpha, num\_iters)

**# Plot the convergence graph**

pyplot.plot(np.arange(len(J\_history)), J\_history, lw=2)

pyplot.xlabel('Number of iterations')

pyplot.ylabel('Cost J')

**# Display the gradient descent's result**

print('theta computed from gradient descent: {}'.format(str(theta)))

price = np.dot(np.array([1.,(1650. - mu[0])/sigma[0] , (3. - mu[1])/sigma[1]]),np.array(theta) ) # You should change this

print(price)

# =====

print('Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): \${:.0f}'.format(price))

pyplot.figure(figsize=(15,10))

alpha=[0.01, 0.003,0.001]

for i in range(len(alpha)):

pyplot.subplot(2,3,i+1)

num\_iters = 400



**# init theta and run gradient descent**

```
theta = np.zeros(3)
```

```
theta, J_history = gradientDescentMulti(X, y, theta, alpha[i], num_iters)
```

**# Plot the convergence graph**

```
pyplot.plot(np.arange(len(J_history)), J_history, lw=2)
```

```
pyplot.xlabel('Number of iterations')
```

```
pyplot.ylabel('Cost J')
```

**# Display the gradient descent's result**

```
print('theta computed from gradient descent: {}'.format(str(theta)))
```

**# Estimate the price of a 1650 sq-ft, 3 br house**

**# Recall that the first column of X is all-ones.**

**# Thus, it does not need to be normalized.**

```
price = np.dot(np.array([1.,(1650. - mu[0])/sigma[0] , (3. - mu[1])/sigma[1]]),np.array(theta) ) #  
You should change this
```

```
print(price)
```

```
print('Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): ${:.0f}'.format(price))
```

```
print("\n")
```

```
pyplot.title('Alpha = '+str(alpha[i]))
```

**Résultat**

289221.5473712181


Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): \$289222

theta computed from gradient descent: [238067.09470609 67368.83752161 19746.42023768]  
203553.10370302649

Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): \$203553

theta computed from gradient descent: [112272.89290139 33254.00585532 14524.42608546]  
94156.27135253145

Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): \$94156



Donc Alpha= 0.01 est le taux d'apprentissage dans une bonne fourchette

# Choose some alpha value - change this

```
alpha = [0.1, 0.3, 0.01, 0.03, 0.001, 0.003, 0.0003, 0.0001]
```

```
num_iters = 2000
```

# init theta and run gradient descent

```
for i in range(len(alpha)):
```

```
    theta = np.zeros(3)
```

```
    theta, J_history = gradientDescentMulti(X, y, theta, alpha[i], num_iters)
```

# Plot the convergence graph

```
pyplot.plot(np.arange(len(J_history)), J_history, lw=2, label = 'learningrate'+str(alpha[i]))
```

```
pyplot.xlabel('Number of iterations')
```

```

pyplot.ylabel('Cost J')

pyplot.legend()

print('theta computed with from gradient descent: {}'.format(str(theta)))

# Display the gradient descent's result

# Estimate the price of a 1650 sq-ft, 3 br house

# Recall that the first column of X is all-ones.

# Thus, it does not need to be normalized.

price = 0 # You should change this

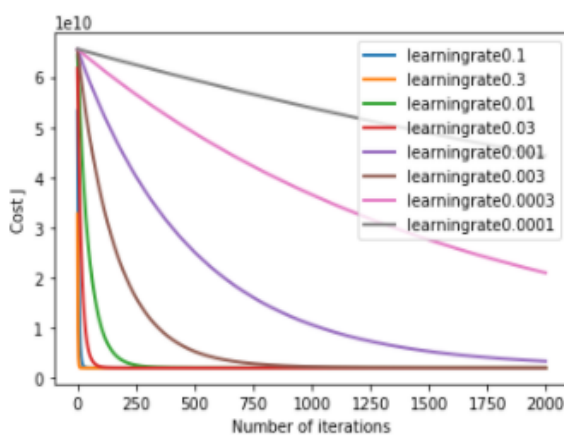
print('Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): ${:.0f}'.format(price))

```

```

theta computed with from gradient descent: [340412.65957447 109447.79646964 -6578.35485416]
theta computed with from gradient descent: [340412.65957447 109447.79646964 -6578.35485416]
theta computed with from gradient descent: [340412.65894002 109439.22578243 -6569.78416695]
theta computed with from gradient descent: [340412.65957447 109447.79646948 -6578.35485399]
theta computed with from gradient descent: [294388.89339564 83125.36792731 15212.40521995]
theta computed with from gradient descent: [339576.43615572 105311.60418477 -2450.82887525]
theta computed with from gradient descent: [153607.04755151 44728.60266957 17804.15924334]
theta computed with from gradient descent: [61709.1336788 18673.5465658 8898.70264165]
Predicted price of a 1650 sq-ft, 3 br house (using gradient descent): $0

```



**Exercise: Ecrivez une methode qui permet de predire le prix d'une maison en fonction de la superficie et du nombre de piece**

```
def predictnorm(superficie, pieces,theta):
```

```
    theta = normalEqn(X,y)
```

```
    price = np.dot(theta.T, np.array([1., superficie, pieces]))
```

```
    return price
```

#### Appel de fonction

```
price = predictnorm(1650,3,theta)
```

```
print(f"une maison avec 500 de superficie et 10 pieces est estimée à {price} en utilisant l'equation normale")
```

#### Résultat

```
une maison avec 500 de superficie et 10 pieces est estimée à 293081.4643350591 en utilisant l'equation normale
```

---

### **4 Exercice: Trouvez des données et faites de la prévision en utilisant les algorithmes implémentées dans ce notebook.**

#### #Importation des librairies

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
%matplotlib inline
```

#### #Chargement des donnees

```
data=pd.read_csv('data/Salaire_employes.csv')
```

```
data
```

	AnnesExperience	Salaire
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0

### #Chargement restreint des données

```
data.head()
```

	AnnesExperience	Salaire
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

---

### #Description des données

```
data.describe()
```

	AnnexExperience	Salaire
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

**#Estimation du salaire en fonction de l'année d' expérience a travers un diagramme**


```
data['Salaire'].value_counts().plot(kind='bar')
```

```
plt.title("Salaire Estimé en fonction de l'Annees d'experiences")
```

```
plt.xlabel('Salaires')
```

```
plt.ylabel('Annees Experiences')
```

```
sns.despine
```



**#Méthode qui nous permet de predire le salaire en fonction de l'annee d'experience**

```
def predict(AnnesExperience, mu, sigma, theta):
```

```
    # init theta and run gradient descent
```

```
    n=X.shape[1]
```

```
    theta = np.zeros(n)
```

```
    theta, J_history = gradientDescentMulti(X, y, theta, alpha=0.003, num_iters=400)
```

```
    Salaire = np.dot(np.array([1.,(AnnesExperience - mu[0])/sigma[0]]),np.array(theta) )
```

```
    return (f"Pour {AnnesExperience} Annees Experiences , le Salaire estimé est %.3f"%Salaire)
```

```
#Exemple deprediction du salaire pour une annee d'experience de 2.2 ans
```

```
salary= predict(2.2,mu,sigma,theta)
```

```
print(salary)
```

```
Pour 2.2 Annees Experiences , le Salaire estimé est 32939.214
```

---